

# Performance with in database JSON



# Contents

<b>PERFORMANCE WITH IN DATABASE JSON .....</b>	<b>1</b>
<b>INITIAL REQUIREMENTS .....</b>	<b>3</b>
<b>JSON DOCUMENTS IN THE DATABASE.....</b>	<b>3</b>
JSON QUERIES .....	3
<i>Access the CDB, PDB creation and administration .....</i>	<i>3</i>
<i>Create a JSON table and populate it.....</i>	<i>4</i>
<i>Use JSON and non JSON tables in the same query .....</i>	<i>5</i>
<i>Create a partitioned JSON table .....</i>	<i>6</i>
<i>Establish a baseline .....</i>	<i>8</i>
<i>Rewrite the query on the JSON table .....</i>	<i>10</i>
<i>Leverage partition pruning .....</i>	<i>12</i>
<i>Create an index on a JSON field .....</i>	<i>14</i>
<i>Partition the table on a JSON attribute.....</i>	<i>19</i>
<i>Create a SEARCH index.....</i>	<i>25</i>
<i>JSON and analytical queries.....</i>	<i>30</i>
<i>Populate In-memory column store.....</i>	<i>32</i>
<i>JSON and materialized views .....</i>	<i>34</i>
JSON DOCUMENTS INGESTION.....	37
<i>Use Fast Ingest tables for JSON documents.....</i>	<i>37</i>



# Initial requirements

- SSH private key to Access the database server in the cloud. This private key is provided along with this manual.
- SSH client app, to login to the database server
- Database server public IP

## JSON documents in the database

### JSON queries

#### Access the CDB, PDB creation and administration

Below are described the first steps to access the container database (CDB) and create a pluggable database (PDB).

Access to the database server as "**opc**" using ssh. Then gain access to the "oracle" user, and use Sql\*Plus to run the following commands:

```
ssh -i id_rsa opc@<db server public ip>

## Gain access to the "oracle" user:

[opc@rdbms21coniaas ~]$ sudo su - oracle

## Use Sql*Plus to show the PDB that will be used for the lab

[oracle@rdbms21coniaas ~]$ sqlplus / as sysdba

SQL*Plus: Release 21.0.0.0.0 - Production on Wed Jan 12 15:44:16 2022
Version 21.3.0.0.0

Copyright (c) 1982, 2021, Oracle. All rights reserved.

Connected to:
Oracle Database 21c Enterprise Edition Release 21.0.0.0.0 - Production
Version 21.3.0.0.0

SQL> show pdbs
```



CON_ID	CON_NAME	OPEN MODE	RESTRICTED
2	PDB\$SEED	READ ONLY	NO
3	ORCLPDB1	MOUNTED	
5	PDBSOE	READ WRITE	NO

SQL> exit

## Create a JSON table and populate it

Connect to the PDBSOE PDB, and run the following SQL statements:

```
$ sqlplus soe/soe@rdbms21coniaas:1521/pdbsoe

## Drop the existing tables

drop table OI_JSON_ORDERS purge;
drop table OI_JSON_ORDER_ITEMS purge;

## Create the new tables
## Regular scalar column ID is used as a primary key
## O_JSON column will be used to store JSON documents, and we enforce a check
constraint on it, to check the correct JSON semantic of the inserted values

create table OI_JSON_ORDERS
(
  ID number(12),
  O_JSON VARCHAR2(4000),
  CONSTRAINT O_JSON_insert_pk primary Key (id),
  CONSTRAINT O_JSON_check CHECK (O_JSON IS JSON)
);

## Now we will create new data in the OI_JSON_ORDERS table
## We use database JSON API to format a JSON document from the ORDERS
relational table

alter session enable parallel DML;
alter session force parallel query parallel 2;
set timing on

insert /*+ APPEND NOLOGGING */
into OI_JSON_ORDERS (id,O_JSON)
select O.order_id,
       json_object (
         'ORDER_ID' value O.ORDER_ID,
         'ORDER_DATE' value O.ORDER_DATE,
         'ORDER_MODE' value O.ORDER_MODE,
         'CUSTOMER_ID' value O.CUSTOMER_ID,
         'ORDER_STATUS' value O.ORDER_STATUS,
         'ORDER_TOTAL' value O.ORDER_TOTAL,
         'SALES_REP_ID' value O.SALES_REP_ID,
```



```

        'PROMOTION_ID' value O.PROMOTION_ID,
        'WAREHOUSE_ID' value O.WAREHOUSE_ID,
        'DELIVERY_TYPE' value O.DELIVERY_TYPE,
        'COST_OF_DELIVERY' value O.COST_OF_DELIVERY,
        'WAIT_TILL_ALL_AVAILABLE' value O.WAIT_TILL_ALL_AVAILABLE,
        'DELIVERY_ADDRESS_ID' value O.DELIVERY_ADDRESS_ID,
        'CUSTOMER_CLASS' value O.CUSTOMER_CLASS,
        'CARD_ID' value O.CARD_ID,
        'INVOICE_ADDRESS_ID' value O.INVOICE_ADDRESS_ID
    ) as mijson
from orders O;

```

1429790 rows created.

Elapsed: 00:00:33.64

SQL> commit;

Commit complete.

Elapsed: 00:00:01.50

We used the JSON\_OBJECT function to format a JSON document out of relational scalar columns.

## Use JSON and non JSON tables in the same query

Now let's use SQL to write a query that access data in both a relational and a JSON table:

## First we will run a "traditional" query, joining two relational tables.

```

select W.WAREHOUSE_NAME, sum(O.ORDER_TOTAL)
from ORDERS O,
WAREHOUSES W
where W.WAREHOUSE_ID = O.WAREHOUSE_ID
and      W.warehouse_name in
('McsRxswjRxXMFDcobjhEIDdEsO', '5eH6XK38SRmNEZCug43EDIjDlCDhbV', 'PLlypy')
group by W.WAREHOUSE_NAME
order by 1
/

```

WAREHOUSE_NAME	SUM(O.ORDER_TOTAL)
5eH6XK38SRmNEZCug43EDIjDlCDhbV	7190505
McsRxswjRxXMFDcobjhEIDdEsO	7368607
PLlypy	7197962

Elapsed: 00:00:00.29

## Now let's write a slightly distinct query to join WAREHOUSES and OI\_JSON\_ORDERS table. We can nicely join the two tables on WAREHOUSE\_ID, which is a regular column in the WAREHOUSES table, and \$.WAREHOUSE\_ID, which is a field of the JSON document stored in OI\_JSON\_ORDERS table.



```

select W.WAREHOUSE_NAME, sum(to_number(json_value (OI.O_JSON,
'$.ORDER_TOTAL')))) as TOTAL
from    OI_JSON_ORDERS OI,
        WAREHOUSES W
where    W.WAREHOUSE_ID = json_value (OI.O_JSON, '$.WAREHOUSE_ID')
and      W.warehouse_name in
('McsRxswjRxXMFDcobjhEIDdEs0', '5eH6XK38SRmNEZCUg43EDIjDICDhbV', 'PLlpy')
group by W.WAREHOUSE_NAME
order by 1;

```

WAREHOUSE_NAME	TOTAL
5eH6XK38SRmNEZCUg43EDIjDICDhbV	7190505
McsRxswjRxXMFDcobjhEIDdEs0	7368607
PLlpy	7197962

Elapsed: 00:00:12.84

The result of the two query is the same, but the performance of the second query is significantly worse than the performance of the 100% relational query.

In the remaining of the workshop, we will demonstrate several ways to speed-up queries involving JSON documents.

## Create a partitioned JSON table

Next we will create and populate partitioned JSON table:

```

## Create an interval partitioned JSON table
## The partition key is a regular column, but could be a JSON field, as we will
see further

create table OI_JSON_ORDER_ITEMS
(
  ID number(12),
  order_date DATE NOT NULL,
  OI_json VARCHAR2(4000),
  CONSTRAINT oi_json_insert_pk primary Key (id),
  CONSTRAINT OI_json_check CHECK (OI_json IS JSON)
)
PARTITION BY RANGE (order_date)
INTERVAL(NUMTOYMINTERVAL(1,'MONTH'))
(
  PARTITION OIJSON_P0 VALUES LESS THAN (TO_DATE('2007-02-01', 'YYYY-MM-DD'))
);

## Disable in-memory for that table, and populate it

alter table OI_JSON_ORDER_ITEMS no inmemory;

```



```

alter session enable parallel DML;
alter session force parallel query parallel 2;
set timing on

insert /*+ APPEND NOLOGGING */
into OI_JSON_ORDER_ITEMS (id,order_date,OI_json)
select O.order_id, O.order_date,
       json_object (
         'ORDER_ID' value O.ORDER_ID,
         'ORDER_DATE' value O.ORDER_DATE,
         'ORDER_MODE' value O.ORDER_MODE,
         'CUSTOMER_ID' value O.CUSTOMER_ID,
         'ORDER_STATUS' value O.ORDER_STATUS,
         'ORDER_TOTAL' value O.ORDER_TOTAL,
         'SALES_REP_ID' value O.SALES_REP_ID,
         'PROMOTION_ID' value O.PROMOTION_ID,
         'WAREHOUSE_ID' value O.WAREHOUSE_ID,
         'DELIVERY_TYPE' value O.DELIVERY_TYPE,
         'COST_OF_DELIVERY' value O.COST_OF_DELIVERY,
         'WAIT_TILL_ALL_AVAILABLE' value O.WAIT_TILL_ALL_AVAILABLE,
         'DELIVERY_ADDRESS_ID' value O.DELIVERY_ADDRESS_ID,
         'CUSTOMER_CLASS' value O.CUSTOMER_CLASS,
         'CARD_ID' value O.CARD_ID,
         'INVOICE_ADDRESS_ID' value O.INVOICE_ADDRESS_ID,
         'ITEMS' value json_arrayagg (
           json_object (
             'ORDER_ID' value OI.ORDER_ID,
             'LINE_ITEM_ID' value OI.LINE_ITEM_ID,
             'PRODUCT_ID' value OI.PRODUCT_ID,
             'UNIT_PRICE' value OI.UNIT_PRICE,
             'QUANTITY' value OI.QUANTITY,
             'DISPATCH_DATE' value OI.DISPATCH_DATE,
             'RETURN_DATE' value OI.RETURN_DATE,
             'GIFT_WRAP' value OI.GIFT_WRAP,
             'CONDITION' value OI.CONDITION,
             'SUPPLIER_ID' value OI.SUPPLIER_ID,
             'ESTIMATED_DELIVERY' value OI.ESTIMATED_DELIVERY
           )
         ) as mijson
from orders O,
     order_items OI
where O.order_id = OI.order_id(+)
group by
O.ORDER_ID,
O.ORDER_DATE,
O.ORDER_MODE,
O.CUSTOMER_ID,
O.ORDER_STATUS,
O.ORDER_TOTAL,
O.SALES_REP_ID,
O.PROMOTION_ID,
O.WAREHOUSE_ID,
O.DELIVERY_TYPE,

```



```
O.COST_OF_DELIVERY,  
O.WAIT_TILL_ALL_AVAILABLE,  
O.DELIVERY_ADDRESS_ID,  
O.CUSTOMER_CLASS,  
O.CARD_ID,  
O.INVOICE_ADDRESS_ID;
```

```
1429790 rows created.
```

```
Elapsed: 00:01:47.06
```

```
SQL> SQL> commit;
```

```
Commit complete.
```

```
Elapsed: 00:00:00.13
```

```
## Exit the session
```

```
exit
```

Now we are ready to start with the performance tests.

## Establish a baseline

We are going to retrieve the sum of unit\_price, for customer\_id= 733116 and date 24-FEB-09. Our baseline is the response time of that query against the relational tables:

```
sqlplus soe/soe@rdbms21coniaas:1521/pdbsoe
```

```
set timing on
```

```
select sum(OI.unit_price)  
from orders O, order_items OI  
where O.customer_id = 733116  
and trunc(O.order_date) = to_date('24-FEB-09','DD-MON-RR')  
and O.order_id = OI.order_id;
```

```
SUM(OI.UNIT_PRICE)
```

```
-----  
5146
```

```
Elapsed: 00:00:01.16
```

```
## Now let's capture the execution plan:
```

```
set autotrace traceonly explain statistics
```

```
select sum(OI.unit_price)  
from orders O, order_items OI
```





```

where O.customer_id = 733116
and trunc(O.order_date) = to_date('24-FEB-09','DD-MON-RR')
and O.order_id = OI.order_id;

```

Elapsed: 00:00:00.01

Execution Plan

Plan hash value: 2506602772

```

-----
-
-----
| Id | Operation                               | Name           | Rows  | Bytes |
|----|-----|-----|-----|-----|
|----|-----|-----|-----|-----|
-
-----
|  0 | SELECT STATEMENT                       |                |      1 |    32 |
|  9 |   (0)| 00:00:01 |                |      |
|  1 |    SORT AGGREGATE                      |                |      1 |    32 |
|  2 |      NESTED LOOPS                      |                |      1 |    32 |
|  9 |        (0)| 00:00:01 |                |      |
|  3 |        NESTED LOOPS                    |                |      3 |    32 |
|  9 |          (0)| 00:00:01 |                |      |
|*  4 |          TABLE ACCESS BY INDEX ROWID BATCHED | ORDERS         |      1 |    22 |
|  5 |            (0)| 00:00:01 |                |      |
|*  5 |            INDEX RANGE SCAN              | ORD_CUSTOMER_IX |      2 |      |
|  3 |              (0)| 00:00:01 |                |      |
|*  6 |            INDEX RANGE SCAN              | ITEM_ORDER_IX  |      3 |      |
|  2 |              (0)| 00:00:01 |                |      |
|  7 |          TABLE ACCESS BY INDEX ROWID   | ORDER_ITEMS    |      3 |      |
30 |  4 |            (0)| 00:00:01 |                |      |
-----
-
-----

```

Predicate Information (identified by operation id):

-----



```
4 - filter(TRUNC(INTERNAL_FUNCTION("O"."ORDER_DATE"))=TO_DATE('24-FEB-09','DD-MON-RR'))
```

```
5 - access("O"."CUSTOMER_ID`=733116)
6 - access("O"."ORDER_ID`= "OI"."ORDER_ID")
```

Note

-----

- this is an adaptive plan

Statistics

-----

```
0 recursive calls
0 db block gets
10 consistent gets
0 physical reads
0 redo size
579 bytes sent via SQL*Net to client
52 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed
```

The execution plan uses a couple of indexes, one indexing ORDER.CUSTOMER\_ID and one indexing ORDER\_ITEMS.ITEM\_ID.

## Rewrite the query on the JSON table

In the next step, we rewrite the query to access OI\_JSON\_ORDER\_ITEMS instead of the two relational tables, ORDERS and ORDER\_ITEMS:

```
set autotrace off

select sum(ARR.UNIT_PRICE)
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
                                     '$.UNIT_PRICE')))
     ) as ARR
where json_value (OIJ.OI_JSON, '$.CUSTOMER_ID') = 733116
and trunc(ARR.ORDER_DATE) = to_date('24-FEB-09','DD-MON-RR');

SUM(ARR.UNIT_PRICE)
```



```

-----
5146
Elapsed: 00:00:03.20

```

Pay attention to and review the syntax: we are using native JSON API to access UNIT\_PRICE in a nested path.

As expected, the result is the same. But the response time is three times higher. Let's capture the execution plan:

```

set autotrace traceonly explain statistics

select sum(ARR.UNIT_PRICE)
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
                             '$.UNIT_PRICE')))
      ) as ARR
where json_value (OIJ.OI_JSON, '$.CUSTOMER_ID') = 733116
and trunc(ARR.ORDER_DATE) = to_date('24-FEB-09','DD-MON-RR');

Elapsed: 00:00:03.25

Execution Plan
-----
Plan hash value: 1435571113

-----
-
-----

| Id | Operation                | Name                | Rows  | Bytes | Cost (%C
PU)| Time      | Pstart| Pstop |
-----
-
-----

|  0 | SELECT STATEMENT         |                     |      1 | 1124 | 468K
(3)| 00:00:19 |      |      |
|  1 | SORT AGGREGATE           |                     |      1 | 1124 |
|      |      |      |      |
|  2 | NESTED LOOPS             |                     |      1 | 1167K | 1251M | 468K
(3)| 00:00:19 |      |      |
|  3 | PARTITION RANGE ALL      |                     | 14298 | 15M | 69606
(1)| 00:00:03 |      1 |1048575|

```



```
|* 4 |      TABLE ACCESS FULL      | OI_JSON_ORDER_ITEMS | 14298 |      15M| 69606
(1)| 00:00:03 |      1 |1048575|
```

```
|* 5 |      JSONTABLE EVALUATION      |      |      |
|      |      |      |
```

```
-----
-
-----

Predicate Information (identified by operation id):
-----
```

```
4 - filter(TO_NUMBER(JSON_VALUE("OIJ"."OI_JSON" FORMAT JSON ,
'$.CUSTOMER_ID'
RETURNING
```

```
      VARCHAR2(4000) NULL ON ERROR))=733116)
```

```
5 - filter(TRUNC("P"."ORDER_DATE")=TO_DATE('24-FEB-09','DD-MON-RR'))
```

```
Statistics
-----
```

```
44 recursive calls
0 db block gets
252515 consistent gets
252160 physical reads
0 redo size
580 bytes sent via SQL*Net to client
52 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
3 sorts (memory)
0 sorts (disk)
1 rows processed
```

The response time is explained by the "full table scan" access to OI\_JSON\_ORDER\_ITEMS table: this is due to the lack of indexes on that table. Note that the response time is anyway pretty acceptable, as the table contains nearly 1.5 million rows.

In the following steps, we will improve that response time.

## Leverage partition pruning

In the previous execution, partition pruning does not occur, because we are using the predicate " and trunc(ARR.ORDER\_DATE) = to\_date('24-FEB-09','DD-MON-RR')". As ARR.ORDER\_DATE is a field in the JSON document, and not the regular column OI\_JSON\_ORDER\_ITEMS.ORDER\_DATE, the predicate is applied as a filter, but not an access path.

Writing the query differently will leverage partition pruning and boost the performance:



```
-- Use ORDER_DATE regular column instead of ORDER_DATE in the JSON document

sqlplus soe/soe@rdbms21coniaas:1521/pdbsoe

set autotrace traceonly explain statistics

select sum(ARR.UNIT_PRICE)
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
                                     '$.UNIT_PRICE')))
     ) as ARR
where json_value (OIJ.OI_JSON, '$.CUSTOMER_ID') = 733116
and OIJ.ORDER_DATE between to_date('24-FEB-09 00:00:00', 'DD-MON-RR HH24:MI:SS')
and to_date('24-FEB-09 23:59:59', 'DD-MON-RR HH24:MI:SS');
```

Elapsed: 00:00:00.05

Execution Plan

Plan hash value: 229940177

```
-----
-
-----
```

	Id	Operation	Name	Rows	Bytes	Cost
	(%CPU)	Time	Pstart Pstop			
	0	SELECT STATEMENT		1	1130	6982
5	(1)	00:00:03				
	1	SORT AGGREGATE		1	1130	
*	2	FILTER				
	3	NESTED LOOPS		64991	70M	6982
5	(1)	00:00:03				
	4	PARTITION RANGE ITERATOR		8	9024	6960
6	(1)	00:00:03	KEY KEY			



* 5	TABLE ACCESS FULL	OI_JSON_ORDER_ITEMS	8	9024
6960				
6 (1)	00:00:03	KEY   KEY		
6	JSONTABLE EVALUATION			
-----				
-				
-----				
Predicate Information (identified by operation id):				
-----				
2 - filter(TO_DATE('24-FEB-09 23:59:59','DD-MON-RR				
HH24:MI:SS')>=TO_DATE('24-				
FEB-09				
00:00:00','DD-MON-RR HH24:MI:SS'))				
5 - filter(TO_NUMBER(JSON_VALUE("OIJ"."OI_JSON" FORMAT JSON ,				
'\$.CUSTOMER_ID'				
RETURNING VARCHAR2(4000)				
NULL ON ERROR))=733116 AND "OIJ"."ORDER_DATE"<=TO_DATE('24-FEB-09				
23:59:59','DD-MON-RR HH24:MI:SS') AND				
"OIJ"."ORDER_DATE">=TO_DATE('24-FEB-09 00:00:00','DD-MON-RR HH24:M				
I:SS'))				
Statistics				
-----				
21 recursive calls				
0 db block gets				
3660 consistent gets				
3646 physical reads				
0 redo size				
580 bytes sent via SQL*Net to client				
52 bytes received via SQL*Net from client				
2 SQL*Net roundtrips to/from client				
0 sorts (memory)				
0 sorts (disk)				
1 rows processed				

The access path is still FULL TABLE SCAN, but after pruning to the corresponding partition.

Create an index on a JSON field



To boost the query even more, we can create an index on a JSON field. We will then benefit from both partitioning and indexing.

```
-- We can index a json field to speed-up the query
-- We would typically index CUSTOMER_ID
-- Pay attention to the syntax, we use native JSON API to define the index !!!

create index I_CUST_ID on
OI_JSON_ORDER_ITEMS
(
    json_value (OI_JSON, '$.CUSTOMER_ID' returning NUMBER(12) error on error
null on empty)
) LOCAL;

Index created.

Elapsed: 00:00:05.14

## Collect standard optimizer statistics on the index

exec dbms_stats.gather_index_stats ('SOE','I_CUST_ID')

select sum(ARR.UNIT_PRICE)
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
'$.UNIT_PRICE')))
        ) as ARR
where json_value (OI_JSON, '$.CUSTOMER_ID' returning NUMBER(12) error on error
null on empty) = 733116
and trunc(ARR.ORDER_DATE) = to_date('24-FEB-09','DD-MON-RR');

SUM(ARR.UNIT_PRICE)
-----
                5146

Elapsed: 00:00:00.01

## The response time was boosted, let's review the execution plan

set autotrace traceonly explain statistics

select sum(ARR.UNIT_PRICE)
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
'$.UNIT_PRICE')))
        ) as ARR
where json_value (OI_JSON, '$.CUSTOMER_ID' returning NUMBER(12) error on error
null on empty) = 733116
```



```
and trunc(ARR.ORDER_DATE) = to_date('24-FEB-09','DD-MON-RR');
```

#### Execution Plan

Plan hash value: 1296896600

Id	Operation	Name	Row			
s	Bytes	Cost (%CPU)	Time	Pstart	Pstop	
0	SELECT STATEMENT					
1	1136	404K (3)	00:00:16			
1	1136					
2	NESTED LOOPS					11
67K	1265M	404K (3)	00:00:16			
3	PARTITION RANGE ALL					142
98	15M	5785 (1)	00:00:01	1	1048575	
4	TABLE ACCESS BY LOCAL INDEX ROWID BATCHED	OI_JSON_ORDER_ITEMS				
142						
98	15M	5785 (1)	00:00:01	1	1048575	
* 5	INDEX RANGE SCAN	I_CUST_ID				57
19	65 (0)	00:00:01		1	1048575	
* 6	JSONTABLE EVALUATION					

#### Predicate Information (identified by operation id):

5 - access(JSON\_VALUE("OI\_JSON" FORMAT JSON , '\$.CUSTOMER\_ID' RETURNING  
NUMBE  
R(12,0) ERROR ON ERROR NULL ON  
EMPTY)=733116)  
6 - filter(TRUNC("P"."ORDER\_DATE")=TO\_DATE('24-FEB-09','DD-MON-RR'))





## Statistics

```

-----
 61 recursive calls
  0 db block gets
194 consistent gets
  2 physical reads
  0 redo size
580 bytes sent via SQL*Net to client
 52 bytes received via SQL*Net from client
  2 SQL*Net roundtrips to/from client
  0 sorts (memory)
  0 sorts (disk)
  1 rows processed

```

Partition pruning does not occur (PARTITION RANGE ALL), but now the index is used.

We can manage to use both partition pruning and indexing to boost even more the performance.

```
set autotrace traceonly explain statistics
```

```

select sum(ARR.UNIT_PRICE)
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
                                     '$.UNIT_PRICE'))
                ) as ARR
where json_value (OI_JSON, '$.CUSTOMER_ID' returning NUMBER(12) error on error
null on empty) = 733116
and OIJ.ORDER_DATE between to_date('24-FEB-09 00:00:00','DD-MON-RR HH24:MI:SS')
and to_date('24-FEB-09 23:59:59','DD-MON-RR HH24:MI:SS');

```

## Execution Plan

```
-----
Plan hash value: 2462359892
```

```

-----
-
-----
| Id | Operation | Name | | | |
|---|---|---|---|---|---|
| Rows | Bytes | Cost (%CPU)| Time | Pstart| Pstop |
|-----|-----|-----|
-
-----
| 0 | SELECT STATEMENT | |
| 1 | 1134 | 32 (0)| 00:00:01 | | |

```



1	SORT AGGREGATE						
	1	1134					
* 2	FILTER						
3	NESTED LOOPS						
	319	353K	32	(0)	00:00:01		
4	PARTITION RANGE AND						
	1	1132	3	(0)	00:00:01	KEY(AP)	KEY(AP)
* 5	TABLE ACCESS BY LOCAL INDEX ROWID BATCHED WITH ZONEMAP						
	OI_JSON_ORDER_ITEMS	1	1132	3	(0)	00:00:01	KEY(AP) KEY(AP)
* 6	INDEX RANGE SCAN						I_CUST_ID
	1		1	(0)	00:00:01	KEY(AP)	KEY(AP)
7	JSONTABLE EVALUATION						

Predicate Information (identified by operation id):

2 - filter(TO\_DATE('24-FEB-09 23:59:59','DD-MON-RR HH24:MI:SS')>=TO\_DATE('24-FEB-09 00:00:00','DD-MON-RR HH24:MI:SS'))

5 - filter(SYS\_ZMAP\_FILTER('/ \* ZM\_PRUNING \*/ SELECT zm."ZONE\_ID\$", CASE WHEN BITAND(zm."ZONE\_STATE\$",1)=1 THEN 1 ELSE CASE WHEN

(zm."MAX\_1\_ORDER\_DATE" < :1 OR zm."MIN\_1\_ORDER\_DATE" > :2) THEN 3 ELSE 2 END END FROM "SOE"."ZMAP\$\_OI\_JSON\_ORDER\_ITEMS" zm WHERE

zm."ZONE\_LEVEL\$ "=0 ORDER BY zm."ZONE\_ID\$",SYS\_OP\_ZONE\_ID(ROWID),TO\_DATE('24-FEB-09 00:00:00','DD-MON-RR HH24:MI:SS'),TO\_DATE('24-FEB-09

23:59:59','DD-MON-RR HH24:MI:SS'))<3 AND "OIJ"."ORDER\_DATE">=TO\_DATE('24-FEB-09 00:00:00','DD-MON-RR HH24:MI:SS') AND

"OIJ"."ORDER\_DATE"<=TO\_DATE('24-FEB-09 23:59:59','DD-MON-RR HH24:MI:SS'))

6 - access(JSON\_VALUE("OI\_JSON" FORMAT JSON , '\$.CUSTOMER\_ID' RETURNING NUMBER(12,0) ERROR ON ERROR NULL ON EMPTY)=733116)



#### Statistics

```
-----
440 recursive calls
5 db block gets
428 consistent gets
29 physical reads
876 redo size
580 bytes sent via SQL*Net to client
52 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
35 sorts (memory)
0 sorts (disk)
1 rows processed
```

We observe that partition pruning and indexing work together, and boost even more the performance.

Let's build a table to compare the different results:

Test case	Response time	Buffer Gets
Relational query	1,16 s	10
Query on JSON table	3,20 s	252515
JSON table + partition pruning	0,05 s	3660
JSON table + Index on json attribute customer_id	0,01 s	194
JSON table + Index on json attribute customer_id + partition pruning	0,01 s	428

We can speed up queries on the JSON table by using traditional tuning techniques, like indexing or partitioning. This demonstrates that all the performance feature of the Oracle database still apply when using JSON data.

In the following steps, we will go further with indexing, but now let's examine another way to partition a table, based on a JSON attribute.

### Partition the table on a JSON attribute

Instead of partitioning the table on a relational scalar column, we could partition on on JSON attribute, and observe the same partition pruning mechanism.

Connect to soe schema and create a JSON table, using a JSON attribute as the partition key:

```
sqlplus soe/soe@rdbms21coniaas:1521/pdbsoe
```

```
CREATE TABLE OI_JSON_ORDER_ITEMS_PART
(id NUMBER(12) NOT NULL PRIMARY KEY,
 OI_JSON VARCHAR2(4000),
```



```

ORDER_DATE DATE GENERATED ALWAYS AS
(json_value (OI_JSON, '$.ORDER_DATE' RETURNING DATE))
)
PARTITION BY RANGE (ORDER_DATE) INTERVAL(NUMTOYMINTERVAL(1, 'MONTH'))
(
PARTITION OIJSONPART_P0 VALUES LESS THAN (TO_DATE('2007-02-01', 'YYYY-MM-
DD'))
);

```

desc OI\_JSON\_ORDER\_ITEMS\_PART

Name	Null?	Type
-----	-----	-----
ID	NOT NULL	NUMBER(12)
OI_JSON		VARCHAR2(4000)
ORDER_DATE		DATE

We create the ORDER\_DATE column as a virtual column on top of the ORDER\_DATE attribute onto the JSON document, then we use that virtual column as the partition key, and interval partition the table with it. We observe, with a "desc" command, that the virtual appears as a regular column.

We will now populate this new table:

```

alter table OI_JSON_ORDER_ITEMS_PART no inmemory;

alter session enable parallel DML;
alter session force parallel query parallel 2;
set timing on

insert /*+ APPEND NOLOGGING */
into OI_JSON_ORDER_ITEMS_PART (id,OI_json)
select O.order_id,
       json_object (
         'ORDER_ID' value O.ORDER_ID,
         'ORDER_DATE' value O.ORDER_DATE,
         'ORDER_MODE' value O.ORDER_MODE,
         'CUSTOMER_ID' value O.CUSTOMER_ID,
         'ORDER_STATUS' value O.ORDER_STATUS,
         'ORDER_TOTAL' value O.ORDER_TOTAL,
         'SALES_REP_ID' value O.SALES_REP_ID,
         'PROMOTION_ID' value O.PROMOTION_ID,
         'WAREHOUSE_ID' value O.WAREHOUSE_ID,
         'DELIVERY_TYPE' value O.DELIVERY_TYPE,
         'COST_OF_DELIVERY' value O.COST_OF_DELIVERY,
         'WAIT_TILL_ALL_AVAILABLE' value O.WAIT_TILL_ALL_AVAILABLE,
         'DELIVERY_ADDRESS_ID' value O.DELIVERY_ADDRESS_ID,
         'CUSTOMER_CLASS' value O.CUSTOMER_CLASS,
         'CARD_ID' value O.CARD_ID,
         'INVOICE_ADDRESS_ID' value O.INVOICE_ADDRESS_ID,
         'ITEMS' value json_arrayagg (
           json_object (
             'ORDER_ID' value OI.ORDER_ID,

```



```

        'LINE_ITEM_ID' value OI.LINE_ITEM_ID,
        'PRODUCT_ID' value OI.PRODUCT_ID,
        'UNIT_PRICE' value OI.UNIT_PRICE,
        'QUANTITY' value OI.QUANTITY,
        'DISPATCH_DATE' value OI.DISPATCH_DATE,
        'RETURN_DATE' value OI.RETURN_DATE,
        'GIFT_WRAP' value OI.GIFT_WRAP,
        'CONDITION' value OI.CONDITION,
        'SUPPLIER_ID' value OI.SUPPLIER_ID,
        'ESTIMATED_DELIVERY' value OI.ESTIMATED_DELIVERY
    )
    )
    ) as mijson
from orders O,
    order_items OI
where O.order_id = OI.order_id(+)
group by
O.ORDER_ID,
O.ORDER_DATE,
O.ORDER_MODE,
O.CUSTOMER_ID,
O.ORDER_STATUS,
O.ORDER_TOTAL,
O.SALES_REP_ID,
O.PROMOTION_ID,
O.WAREHOUSE_ID,
O.DELIVERY_TYPE,
O.COST_OF_DELIVERY,
O.WAIT_TILL_ALL_AVAILABLE,
O.DELIVERY_ADDRESS_ID,
O.CUSTOMER_CLASS,
O.CARD_ID,
O.INVOICE_ADDRESS_ID;

1429790 rows created.

Elapsed: 00:01:44.14
SQL> SQL> SQL> SQL> commit;

Commit complete.

Elapsed: 00:00:00.07

```

Now we will test the same query, to observe partition pruning:

```

-- Exit and reconnect to avoid parallel query execution !!!

exit
sqlplus soe/soe@rdbms21coniaas:1521/pdbsoe
set timing on

select sum(ARR.UNIT_PRICE)

```



```

from OI_JSON_ORDER_ITEMS_PART OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
                                     '$.UNIT_PRICE')))
     ) as ARR
where json_value (OIJ.OI_JSON, '$.CUSTOMER_ID') = 733116
and trunc(ARR.ORDER_DATE) = to_date('24-FEB-09','DD-MON-RR');

```

```

SUM(ARR.UNIT_PRICE)
-----
                5146

```

Elapsed: 00:00:04.70

```
set autotrace traceonly explain statistics
```

```

select sum(ARR.UNIT_PRICE)
from OI_JSON_ORDER_ITEMS_PART OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
                                     '$.UNIT_PRICE')))
     ) as ARR
where json_value (OIJ.OI_JSON, '$.CUSTOMER_ID') = 733116
and trunc(ARR.ORDER_DATE) = to_date('24-FEB-09','DD-MON-RR');

```

Elapsed: 00:00:04.68

Execution Plan

-----  
Plan hash value: 939613720

```

-----
-
-----

```

Id	Operation	Name	Rows	Bytes	Cos
t (%CPU)	Time	Pstart  Pstop			
0	SELECT STATEMENT		1	2006	4
92K (3)	00:00:20				
1	SORT AGGREGATE		1	2006	
2	NESTED LOOPS			1239K	2372M
92K (3)	00:00:20				4



```

| 3 | PARTITION RANGE ALL | 15180 | 28M | 697
81 (1) | 00:00:03 | 1 | 1048575 |
|* 4 | TABLE ACCESS FULL | OI_JSON_ORDER_ITEMS_PART | 15180 | 28M |
697
81 (1) | 00:00:03 | 1 | 1048575 |
|* 5 | JSONTABLE EVALUATION | | | |
| | | | |

```

-----  
 -  
 -----  
 Predicate Information (identified by operation id):  
 -----

```

 4 - filter(TO_NUMBER(JSON_VALUE("OIJ"."OI_JSON" FORMAT JSON ,
'$.CUSTOMER_ID'
RETURNING VARCHAR2(4000)

NULL ON ERROR))=733116)
 5 - filter(TRUNC("P"."ORDER_DATE")=TO_DATE('24-FEB-09','DD-MON-RR'))

```

Note

```

-----
- dynamic statistics used: dynamic sampling (level=2)

```

Statistics

```

-----
      0 recursive calls
      0 db block gets
253590 consistent gets
249728 physical reads
      0 redo size
 580 bytes sent via SQL*Net to client
  52 bytes received via SQL*Net from client
    2 SQL*Net roundtrips to/from client
      0 sorts (memory)
      0 sorts (disk)
      1 rows processed

```

Partition pruning does not kick in, because we are not using the virtual column in the predicate. Change the query to use the virtual column:

```

--- WE should use ORDER_DATE virtual column to leverage partition pruning !!!

select sum(ARR.UNIT_PRICE)

```



```

from OI_JSON_ORDER_ITEMS_PART OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
                                     '$.UNIT_PRICE'))
                ) as ARR
where json_value (OIJ.OI_JSON, '$.CUSTOMER_ID') = 733116
and OIJ.ORDER_DATE between to_date('24-FEB-09 00:00:00','DD-MON-RR HH24:MI:SS')
and to_date('24-FEB-09 23:59:59','DD-MON-RR HH24:MI:SS');

```

Elapsed: 00:00:00.06

#### Execution Plan

Plan hash value: 3037461420

-----						
-						
-----						
Id	Operation	Name		Rows	Bytes	
Cost (%CPU)	Time	Pstart	Pstop			
-----						
-						
-----						
0	SELECT STATEMENT			1	2013	
3935 (54)	00:00:01					
1	SORT AGGREGATE			1	2013	
* 2	FILTER					
3	NESTED LOOPS			250K	480M	
3935 (54)	00:00:01					
4	PARTITION RANGE ITERATOR			31	62341	
3091 (68)	00:00:01	KEY	KEY			
* 5	TABLE ACCESS FULL	OI_JSON_ORDER_ITEMS_PART		31	62341	
3091 (68)	00:00:01	KEY	KEY			
6	JSONTABLE EVALUATION					
-----						
-						
-----						





Predicate Information (identified by operation id):

```
-----  
  
2 - filter(TO_DATE('24-FEB-09 23:59:59','DD-MON-RR  
HH24:MI:SS')>=TO_DATE('24-  
FEB-09 00:00:00','DD-MON-RR  
  
HH24:MI:SS'))  
5 - filter(TO_NUMBER(JSON_VALUE("OIJ"."OI_JSON" FORMAT JSON ,  
'$.CUSTOMER_ID'  
RETURNING VARCHAR2(4000) NULL  
  
ON ERROR))=733116 AND "OIJ"."ORDER_DATE">=TO_DATE('24-FEB-09 00:00  
:00','DD-MON-RR HH24:MI:SS') AND  
  
"OIJ"."ORDER_DATE"<=TO_DATE('24-FEB-09 23:59:59','DD-MON-RR HH24:M  
I:SS'))
```

Note

-----  
- dynamic statistics used: dynamic sampling (level=2)

Statistics

```
-----  
18 recursive calls  
0 db block gets  
3663 consistent gets  
0 physical reads  
0 redo size  
580 bytes sent via SQL*Net to client  
52 bytes received via SQL*Net from client  
2 SQL*Net roundtrips to/from client  
0 sorts (memory)  
0 sorts (disk)  
1 rows processed
```

This illustrates that partitioning might be implemented using a JSON attribute as the partition key.

## Create a SEARCH index

Indexing customer\_id attribute of the JSON column had a great performance impact on the previous query. But what if we don't use customer\_id anymore in the predicate ? What if we cannot anticipate which of the attributes will be used in the predicate ? This is a common case, that can be addressed with a powerful SEARCH index.

```
## Instead of customer_id, we will use card_id in the predicate
```



```

select sum(ARR.UNIT_PRICE)
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
'$$.UNIT_PRICE')))
                ) as ARR
where json_value (OI_JSON, '$.CARD_ID' returning NUMBER(12) error on error
null on empty) = 1465982
and trunc(ARR.ORDER_DATE) = to_date('24-FEB-09','DD-MON-RR');

```

```

SUM(ARR.UNIT_PRICE)
-----
                5146

```

Elapsed: 00:00:07.22

```
set autotrace traceonly explain statistics
```

```

select sum(ARR.UNIT_PRICE)
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path
'$$.UNIT_PRICE')))
                ) as ARR
where json_value (OI_JSON, '$.CARD_ID' returning NUMBER(12) error on error
null on empty) = 1465982
and trunc(ARR.ORDER_DATE) = to_date('24-FEB-09','DD-MON-RR');

```

Elapsed: 00:00:06.28

Execution Plan

Plan hash value: 1435571113

-----						
-						
-----						
Id	Operation	Name	Rows	Bytes	Cost (%C	
PU)	Time	Pstart  Pstop				
-----						
-						
-----						
0	SELECT STATEMENT		1	1124	468K	
(3)	00:00:19					
1	SORT AGGREGATE		1	1124		



	2	NESTED LOOPS		1167K	1251M	468K
(3)	00:00:19					
	3	PARTITION RANGE ALL		14298	15M	69598
(1)	00:00:03	1	1048575			
* (1)	4	TABLE ACCESS FULL	OI_JSON_ORDER_ITEMS	14298	15M	69598
	00:00:03	1	1048575			
* (1)	5	JSONTABLE EVALUATION				

Predicate Information (identified by operation id):

```

4 - filter(JSON_VALUE("OI_JSON" FORMAT JSON , '$.CARD_ID' RETURNING
NUMBER(12
,0) ERROR ON ERROR

        NULL ON EMPTY)=1465982)
5 - filter(TRUNC("P"."ORDER_DATE")=TO_DATE('24-FEB-09','DD-MON-RR'))

```

#### Statistics

```

-----
66 recursive calls
0 db block gets
252357 consistent gets
252160 physical reads
0 redo size
580 bytes sent via SQL*Net to client
52 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed

```

As card\_id attribute is not indexed, we are back to the full table scan execution plan. Let's build a search index on the JSON column (allow something like 10 minutes to complete):

```
create search index I_JSON_SEARCH on OI_JSON_ORDER_ITEMS(OI_JSON) for JSON;
```

Index created.

Elapsed: 00:10:09.64



Now repeat the query and observe the performance boost:

```
set autotrace traceonly explain statistics
```

```
select sum(ARR.UNIT_PRICE)
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                              NESTED PATH '$.ITEMS[*]'
                              COLUMNS (UNIT_PRICE NUMBER path
                              '$.UNIT_PRICE'))
                ) as ARR
where json_value (OI_JSON, '$.CARD_ID') = 1465982
and trunc(ARR.ORDER_DATE) = to_date('24-FEB-09','DD-MON-RR');
```

Elapsed: 00:00:00.03

Execution Plan

Plan hash value: 1938435827

Id	Operation	Name	Rows	Byt
es	Cost (%CPU)   Time	Pstart   Pstop		
0	SELECT STATEMENT		1	11
36	521 (1)   00:00:01			
1	SORT AGGREGATE		1	11
36				
2	NESTED LOOPS		584	6
47K	521 (1)   00:00:01			
* 3	TABLE ACCESS BY GLOBAL INDEX ROWID	OI_JSON_ORDER_ITEMS	7	
79				
24	325 (0)   00:00:01	ROWID   ROWID		
* 4	DOMAIN INDEX	I_JSON_SEARCH		
4	(0)   00:00:01			
* 5	JSONTABLE EVALUATION			



-----  
-  
-----  
  
Predicate Information (identified by operation id):  
-----

```
3 - filter(TO_NUMBER(JSON_VALUE("OI_JSON" FORMAT JSON , '$.CARD_ID'  
RETURNING  
VARCHAR2(4000) NULL ON  
  
ERROR))=1465982)  
4 -  
access("CTXSYS"."CONTAINS"("OIJ"."OI_JSON",'(sdata(FNUM_14173D25B4DD102AB  
9B6F2851AEE2420_CARD_ID = 1465982  
  
))')>0)  
5 - filter(TRUNC("P"."ORDER_DATE")=TO_DATE('24-FEB-09','DD-MON-RR'))
```

Statistics  
-----

```
216 recursive calls  
5 db block gets  
176 consistent gets  
0 physical reads  
1112 redo size  
580 bytes sent via SQL*Net to client  
52 bytes received via SQL*Net from client  
2 SQL*Net roundtrips to/from client  
2 sorts (memory)  
0 sorts (disk)  
1 rows processed
```

Observe that the SEARCH index (DOMAIN INDEX) is being used, and that the performance has been significantly improved.

Note that we could also use the dotted notation to write the query:

```
select sum(ARR.UNIT_PRICE)  
from OI_JSON_ORDER_ITEMS OIJ,  
     json_table(OIJ.OI_JSON,  
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',  
                             NESTED PATH '$.ITEMS[*]'  
                             COLUMNS (UNIT_PRICE NUMBER path  
                             '$.UNIT_PRICE'))  
                ) as ARR  
where OIJ.OI_JSON.CARD_ID = 1465982  
and trunc(ARR.ORDER_DATE) = to_date('24-FEB-09','DD-MON-RR');
```



## JSON and analytical queries

We can run analytical queries on top of a JSON column: connect to soe schema and run the following query:

```
sqlplus soe/soe@rdbms21coniaas:1521/pdbsoe
select sum(ARR.UNIT_PRICE*ARR.QUANTITY) as "GrantTotal"
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path '$.UNIT_PRICE',
                                       QUANTITY NUMBER path '$.QUANTITY'))
                ) as ARR
where OIJ.ORDER_DATE between to_date('01-FEB-09 00:00:00','DD-MON-RR
HH24:MI:SS') and to_date('28-FEB-09 23:59:59','DD-MON-RR HH24:MI:SS');

GrantTotal
-----
278913547

Elapsed: 00:00:00.40

set autotrace traceonly explain statistics

select sum(ARR.UNIT_PRICE*ARR.QUANTITY) as "GrantTotal"
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path '$.UNIT_PRICE',
                                       QUANTITY NUMBER path '$.QUANTITY'))
                ) as ARR
where OIJ.ORDER_DATE between to_date('01-FEB-09 00:00:00','DD-MON-RR
HH24:MI:SS') and to_date('28-FEB-09 23:59:59','DD-MON-RR HH24:MI:SS');

Elapsed: 00:00:00.38

Execution Plan
-----
Plan hash value: 2564720920

-----
-
-----
| Id | Operation                                | Name                | Rows  | Bytes |
| Cost (%CPU)| Time                | Pstart| Pstop |
-----
-
-----
```



0	SELECT STATEMENT	1	1132
630K	(1) 00:00:25		
1	SORT AGGREGATE	1	1132
* 2	FILTER		
3	NESTED LOOPS	168M	177
G 630K	(1) 00:00:25		
4	PARTITION RANGE AND	20624	22
M 69608	(1) 00:00:03 KEY(AP) KEY(AP)		
* 5	TABLE ACCESS FULL WITH ZONEMAP	OI_JSON_ORDER_ITEMS	20624
22			
M 69608	(1) 00:00:03 KEY(AP) KEY(AP)		
6	JSONTABLE EVALUATION		

-----  
-  
-----

Predicate Information (identified by operation id):

```

2 - filter(TO_DATE('28-FEB-09 23:59:59','DD-MON-RR
HH24:MI:SS')>=TO_DATE('01-
FEB-09 00:00:00','DD-MON-RR
HH24:MI:SS'))
5 - filter(SYS_ZMAP_FILTER('/ * ZM_PRUNING */ SELECT zm."ZONE_ID$", CASE WHEN
BITAND(zm."ZONE_STATE$",1)=1 THEN
1 ELSE CASE WHEN (zm."MAX_1_ORDER_DATE" < :1 OR zm."MIN_1_ORDER_DA
TE" > :2) THEN 3 ELSE 2 END END FROM
"SOE"."ZMAP$_OI_JSON_ORDER_ITEMS" zm WHERE zm."ZONE_LEVEL$"=0 ORDE
R BY
zm."ZONE_ID$",SYS_OP_ZONE_ID(ROWID),TO_DATE('01-FEB-09 00:00:00',
'DD-MON-RR HH24:MI:SS'),TO_DATE('28-FEB-09
23:59:59','DD-MON-RR HH24:MI:SS'))<3 AND "OIJ"."ORDER_DATE"<=TO_DA
TE('28-FEB-09 23:59:59','DD-MON-RR HH24:MI:SS')
AND "OIJ"."ORDER_DATE">=TO_DATE('01-FEB-09 00:00:00','DD-MON-RR HH
24:MI:SS'))

```



## Statistics

```
-----
17  recursive calls
0   db block gets
3662 consistent gets
3646 physical reads
0   redo size
574  bytes sent via SQL*Net to client
52   bytes received via SQL*Net from client
2    SQL*Net roundtrips to/from client
0    sorts (memory)
0    sorts (disk)
1    rows processed
```

We observe that the query is resolved by partition pruning, with a fast response time. We can use In-memory database with a JSON table, to leverage in-memory analytical queries. In the following steps, we will populate the in-memory column store with a partition, and observe the result:

## Populate In-memory column store

```
--- Use the following query to get the name of the partitions for 2009-02:

select PARTITION_NAME, HIGH_VALUE from user_tab_partitions where table_name =
'OI_JSON_ORDER_ITEMS';

SYS_P2582
TO_DATE(' 2009-03-01 00:00:00', 'SYYYY-MM-DD HH24:MI:SS',
'NLS_CALENDAR=GREGORIA'

-- Place that partition into the IMC: replace the partition name by your
partition name !!!

alter table OI_JSON_ORDER_ITEMS modify partition SYS_P2582 inmemory priority
critical;

Table altered.

Elapsed: 00:00:00.01

-- Re-run the analytical query and observe the execution plan

set autotrace traceonly explain statistics

select sum(ARR.UNIT_PRICE*ARR.QUANTITY) as "GrantTotal"
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]')
```





```

                                COLUMNS (UNIT_PRICE NUMBER path '$.UNIT_PRICE',
                                QUANTITY NUMBER path '$.QUANTITY'))
                                ) as ARR
where OIJ.ORDER_DATE between to_date('01-FEB-09 00:00:00','DD-MON-RR
HH24:MI:SS') and to_date('28-FEB-09 23:59:59','DD-MON-RR HH24:MI:SS');

```

Elapsed: 00:00:00.40

#### Execution Plan

Plan hash value: 2564720920

```

-----
-
-----
| Id | Operation                               | Name           | Rows
|----|-----|-----|-----|-----|-----|
| 0 | SELECT STATEMENT                       |                |
1 | 1132 | 630K (1) | 00:00:25 |      |      |
| 1 | SORT AGGREGATE                         |                |
1 | 1132 |          |          |      |      |
|* 2 | FILTER                                 |                |
|    |          |          |          |      |      |
| 3 | NESTED LOOPS                           |                |      16
8M | 177G | 630K (1) | 00:00:25 |      |      |
| 4 | PARTITION RANGE AND                     |                |      2062
4 | 22M | 69608 (1) | 00:00:03 | KEY(AP) | KEY(AP) |
|* 5 | TABLE ACCESS INMEMORY FULL WITH ZONEMAP | OI_JSON_ORDER_ITEMS |
2062
4 | 22M | 69608 (1) | 00:00:03 | KEY(AP) | KEY(AP) |
| 6 | JSONTABLE EVALUATION                     |                |
|    |          |          |          |      |      |
-----
-
-----

```

Predicate Information (identified by operation id):

-----



```

2 - filter(TO_DATE('28-FEB-09 23:59:59','DD-MON-RR
HH24:MI:SS')>=TO_DATE('01-
FEB-09 00:00:00','DD-MON-RR HH24:MI:SS'))

5 - inmemory("OIJ"."ORDER_DATE"<=TO_DATE('28-FEB-09 23:59:59','DD-MON-RR
HH24
:MI:SS') AND

"OIJ"."ORDER_DATE">=TO_DATE('01-FEB-09 00:00:00','DD-MON-RR HH24:M
I:SS'))

filter(SYS_ZMAP_FILTER('/ * ZM_PRUNING */ SELECT zm."ZONE_ID$", CASE WHEN
BITAND(zm."ZONE_STATE$",1)=1 THEN 1 ELSE

CASE WHEN (zm."MAX_1_ORDER_DATE" < :1 OR zm."MIN_1_ORDER_DATE" > :
2) THEN 3 ELSE 2 END END FROM

"SOE"."ZMAP$_OI_JSON_ORDER_ITEMS" zm WHERE zm."ZONE_LEVEL$"=0 ORDE
R BY

zm."ZONE_ID$",SYS_OP_ZONE_ID(ROWID),TO_DATE('01-FEB-09 00:00:00',
'DD-MON-RR HH24:MI:SS'),TO_DATE('28-FEB-09

23:59:59','DD-MON-RR HH24:MI:SS'))<3 AND "OIJ"."ORDER_DATE"<=TO_DA
TE('28-FEB-09 23:59:59','DD-MON-RR HH24:MI:SS') AND

"OIJ"."ORDER_DATE">=TO_DATE('01-FEB-09 00:00:00','DD-MON-RR HH24:M
I:SS'))

```

#### Statistics

```

-----
83 recursive calls
0 db block gets
3726 consistent gets
3646 physical reads
0 redo size
574 bytes sent via SQL*Net to client
52 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed

```

## JSON and materialized views

We can create materialized views on top of JSON documents. This will dramatically speed-up the analytical queries. Connect to soe schema, and create a materialized view:



```

sqlplus soe/soe@rdbms21coniaas:1521/pdbsoe

create materialized view fast_mv
build immediate
refresh complete on demand
as
select to_char(OIJ.ORDER_DATE,'YYYYMM') as "Month",
sum(ARR.UNIT_PRICE*ARR.QUANTITY) as "GrantTotal"
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path '$.UNIT_PRICE',
                                       QUANTITY NUMBER path '$.QUANTITY'))
                ) as ARR
group by to_char(OIJ.ORDER_DATE,'YYYYMM');

-- Now compare the performance and scalability metrics if you use either the
base table or the materialized view:

set autotrace traceonly explain statistics
set timing on

select to_char(OIJ.ORDER_DATE,'YYYYMM') as "Month",
sum(ARR.UNIT_PRICE*ARR.QUANTITY) as "GrantTotal"
from OI_JSON_ORDER_ITEMS OIJ,
     json_table(OIJ.OI_JSON,
                '$' COLUMNS (ORDER_DATE DATE path '$.ORDER_DATE',
                             NESTED PATH '$.ITEMS[*]'
                             COLUMNS (UNIT_PRICE NUMBER path '$.UNIT_PRICE',
                                       QUANTITY NUMBER path '$.QUANTITY'))
                ) as ARR
group by to_char(OIJ.ORDER_DATE,'YYYYMM');

64 rows selected.

Elapsed: 00:00:38.50

Execution Plan
-----
Plan hash value: 2840653318

-----
-
-----

| Id | Operation | Name | Rows | Bytes | Tem
pSpc| Cost (%CPU)| Time | Pstart| Pstop |
-----
-
-----

```



0	SELECT STATEMENT			46640	50M
	1987M (1)	21:34:11			
1	HASH GROUP BY			46640	50M
12T	1987M (1)	21:34:11			
2	NESTED LOOPS			11G	12T
	38M (1)	00:25:23			
3	PARTITION RANGE ALL			1429K	1538M
	69555 (1)	00:00:03	1	1048575	
4	TABLE ACCESS INMEMORY FULL	OI_JSON_ORDER_ITEMS		1429K	1538M
	69555 (1)	00:00:03	1	1048575	
5	JSONTABLE EVALUATION				

#### Statistics

```

57 recursive calls
5 db block gets
252342 consistent gets
252160 physical reads
1032 redo size
2519 bytes sent via SQL*Net to client
96 bytes received via SQL*Net from client
6 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
64 rows processed

```

This analytical query runs in 38,5 s against the base table. Let's rewrite the query and access to the materialized view:

```

set autotrace traceonly explain statistics
set timing on
select * from fast_mv;

```

64 rows selected.

Elapsed: 00:00:00.01

Execution Plan

Plan hash value: 140868995



```

-
-----

| Id | Operation | Name | Rows | Bytes | Cost (%CPU)| Time |
-----|-----|-----|-----|-----|-----|-----|-----|
| 0 | SELECT STATEMENT | | 64 | 896 | 3 (0)| 0:00:01 |
| 1 | MAT_VIEW ACCESS INMEMORY FULL | FAST_MV | 64 | 896 | 3 (0)| 0:00:01 |
-----

-
-----

Statistics
-----
1 recursive calls
0 db block gets
7 consistent gets
0 physical reads
0 redo size
2519 bytes sent via SQL*Net to client
96 bytes received via SQL*Net from client
6 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
64 rows processed

```

This concludes the JSON queries part of the lab. In the next chapter, we will review ingestion functionalities.

## JSON documents ingestion

## Use Fast Ingest tables for JSON documents

JSON documents are widely used to create a "schema on read" data model. Many IOT devices use JSON format to send the metrics they are collecting. Ingesting these documents in real time can be challenging, depending on the number of IOT sending information concurrently.



In the next steps, we will review a fast ingestion mechanism that can be used for this purpose: Oracle 19c MEMOPTIMIZE FOR WRITE tables. We will compare unitary inserts into a regular table and into a MEMOPTIMIZE FOR WRITE table.

Connect to soe schema, and create a regular JSON table:

```
sqlplus soe/soe@rdbms21coniaas:1521/pdbsoe

create table OI_JSON_REGULAR
(
  ID number(12),
  OI_json VARCHAR2(4000),
  CONSTRAINT oi_json_regular_pk primary Key (id),
  CONSTRAINT OI_json_regular_check CHECK (OI_json IS JSON)
);

-- Create a MEMOPTIMIZE FOR WRITE table with JSON column !!!

create table OI_JSON_MEMOPT4WRITE
(
  ID number(12),
  OI_JSON varchar2(4000),
  CONSTRAINT oi_json_MEMOPT_pk primary Key (id),
  CONSTRAINT OI_json_MEMOPT_check CHECK (OI_json IS JSON)
) segment creation immediate memoptimize for write;

-- Now we create a PL/SQL block that inserts row by row into the regular table:

create or replace procedure PC_INS_REGULAR (p_num_rows IN PLS_INTEGER)
IS
  CURSOR c_oi (p_num IN PLS_INTEGER)
  IS
    select id, OI_json
    from OI_JSON_ORDER_ITEMS
    where rownum <= p_num;
begin
  FOR cur in c_oi (p_num_rows)
  LOOP
    insert into OI_JSON_REGULAR (id,oi_json) values (cur.id,cur.oi_json);
    commit;
  END LOOP;
END;
/

-- Now we create a PL/SQL block that inserts row by row into the MEMOPTIMIZE
FOR WRITE table:

create or replace procedure PC_INS_MEMOPT4WRITE (p_num_rows IN PLS_INTEGER)
IS
  CURSOR c_oi (p_num IN PLS_INTEGER)
  IS
    select id, OI_json
    from OI_JSON_ORDER_ITEMS
```



```

        where rownum <= p_num;
begin
    FOR cur in c_oi (p_num_rows)
    LOOP
        insert /*+ memoptimize_write */ into OI_JSON_MEMOPT4WRITE (id,oi_json)
values (cur.id,cur.oi_json);
        commit;
    END LOOP;
END;
/

```

Pay attention to the syntax details that are specific to memoptimize for write tables (highlighted in green).

Now we will use the created procedures to insert some rows in both the regular and the memoptimize for write tables. We will then compare the results in terms of performance and throughput: let's start with 1000 rows.

```

set timing on

--- 1000 rows !!!
truncate table OI_JSON_REGULAR;
truncate table OI_JSON_MEMOPT4WRITE;

exec PC_INS_REGULAR(1000)

PL/SQL procedure successfully completed.

Elapsed: 00:00:00.23
SQL> exec PC_INS_MEMOPT4WRITE(1000)

PL/SQL procedure successfully completed.

Elapsed: 00:00:18.12

```

The first time we use the memoptimize for write table, a memory allocation is done in the large pool. This explains why the first execution is so slow, but this will occur only once. Let's re-run the second test:

```

truncate table OI_JSON_MEMOPT4WRITE;

Table truncated.

Elapsed: 00:00:00.07
SQL> exec PC_INS_MEMOPT4WRITE(1000)

PL/SQL procedure successfully completed.

Elapsed: 00:00:00.04

```



Even with only 1000 rows inserted, memoptimize for write table is way faster. Let's try with each time more rows and compare:

```
-- 10.000 rows !!!
truncate table OI_JSON_REGULAR;
truncate table OI_JSON_MEMOPT4WRITE;

SQL> exec PC_INS_REGULAR(10000)

PL/SQL procedure successfully completed.

Elapsed: 00:00:05.30
SQL> exec PC_INS_MEMOPT4WRITE(10000)

PL/SQL procedure successfully completed.

Elapsed: 00:00:00.70

--- 100.000 rows !!!

truncate table OI_JSON_REGULAR;
truncate table OI_JSON_MEMOPT4WRITE;

SQL> exec PC_INS_REGULAR(100000)

PL/SQL procedure successfully completed.

Elapsed: 00:00:24.83
SQL> exec PC_INS_MEMOPT4WRITE(100000)

PL/SQL procedure successfully completed.

Elapsed: 00:00:08.00

-- Count the rows in each table and check:

select count(*) from OI_JSON_REGULAR;

COUNT(*)
-----
100000

select count(*) from OI_JSON_MEMOPT4WRITE;

COUNT(*)
-----
99390
```





The count in the memoptimize for write table doesn't match the number of rows inserted (100.000). This is because rows are committed asynchronously in the MEMOPTIMIZE FOR WRITE table.

We might want to use DBMS\_MEMOPTIMIZE.WRITE\_END procedure to force an immediate flush of the large pool to the table, or just wait for the rows to be eventually flushed automatically. This is important to understand, and might be suitable for your IOT business case (or not).

After some seconds, the missing rows are flushed and we can see them in the table:

```
select count(*) from OI_JSON_MEMOPT4WRITE;

COUNT(*)
-----
    100000

--- 1.000.000 rows !!!

SQL> exec PC_INS_REGULAR(1000000)

PL/SQL procedure successfully completed.

Elapsed: 00:04:28.81

SQL> select count(*) from OI_JSON_REGULAR;

COUNT(*)
-----
    1000000

SQL> exec PC_INS_MEMOPT4WRITE(1000000)

PL/SQL procedure successfully completed.

Elapsed: 00:01:29.13

SQL> select count(*) from OI_JSON_MEMOPT4WRITE;

COUNT(*)
-----
    999875

Elapsed: 00:00:00.15
SQL> exec DBMS_MEMOPTIMIZE.WRITE_END

PL/SQL procedure successfully completed.

Elapsed: 00:00:00.04
SQL> select count(*) from OI_JSON_MEMOPT4WRITE;
```



COUNT(\*)

-----

1000000

Elapsed: 00:00:00.03

